

Digital Signal Processing Lab

Take Home Exam - Q1

Saad Zubairi
shz2020

November 8, 2025

Solution

Premise

The basic premise of this question was to create a program that lets us shift the frequency of an input signal by a scaling factor while also preserving the harmonic structure of the signal, and letting the user shift the frequency through a GUI based slider that repeats the audio in real time. The suggested way of approaching this problem is by taking FFT of the input, shifting the frequency spectrим by the scaling factor, and then inversing the FFT.

a) Approach

The approach was simple. Taking hints from [Demo 13](#), [Demo 14](#), and [Demo 56](#), I implemented frequency scaling using FFT processing with overlapping block processing. To ensure smooth transitions between block boundaries and avoid artifacts, the signal is broken down into overlapping blocks with a 50% overlap factor, obtained through trial and error. Each block is processed independently using FFT, frequency scaling via interpolation, and IFFT. The processed blocks are then combined using the overlap-add method.

Detailed steps to solution

- **Blocks Forming:** Before processing, the signal is divided into overlapping blocks using the function `frames_to_process_with_hops()`. The overlap factor is set to 0.5.

```
1 def frames_to_process_with_hops(all_frames, block_size, overlap_factor):
2     frames = []
3     hop_count = math.floor(block_size * (1 - overlap_factor))
4     idx_i = 0
5     idx_l = block_size
6     while idx_l < len(all_frames) - 1:
7         frames.append(all_frames[idx_i:idx_l])
8         idx_i += hop_count
9         idx_l = idx_i + block_size
10    return frames
11
```

Snippet 1: Block formation with overlap

- **FFT and Frequency Scaling:** Each block is processed using the `process_block_fft_scaling()` function. For each block, this function:

1. Computes the real FFT of the input block.
2. Scales the frequency spectrum by the factor α using linear interpolation
3. Computes the inverse real FFT to obtain the time-domain output

The frequency scaling is achieved by mapping each bin k to a destination bin k/α . Of course since k/α may rarely be an integer, linear interpolation is used between bins:

$$\mathbf{y}[k] = (1 - t) \cdot \mathbf{x}[\lfloor k/\alpha \rfloor] + t \cdot \mathbf{x}[\lfloor k/\alpha \rfloor + 1]$$

where $t = k/\alpha - \lfloor k/\alpha \rfloor$ is of course the fractional part.

```

1 def process_block_fft_scaling(input_block, alpha):
2     X = np.fft.rfft(input_block)
3     Y = np.zeros_like(X)
4     for src_ind in range(X.size):
5         dst_ind = src_ind / alpha
6         if dst_ind < X.size - 1:
7             i0 = int(np.floor(dst_ind))
8             i1 = i0 + 1
9             t = dst_ind - i0
10            Y[src_ind] = (1 - t) * X[i0] + t * X[i1]
11    y = np.fft.irfft(Y)
12    return y
13

```

Snippet 2: FFT-based frequency scaling with interpolation

- **Overlap-Add Processing:** For smooth combinig of blocks, the tail of the previous block (which essentially is the overlapping portion) is stored and added to the beginning of the current block with a 0.5 weighting (a sort of averaging) to prevent abrupt artifacts:

```

1 # overlap and add
2 output_block[:len(prev_tail)] = 0.5 * (output_block[:len(prev_tail)] +
3                                         prev_tail)

```

Snippet 3: Overlap-add processing

- **Real-time Processing Loop:** The main processing loop continuously:

1. Reads the current scaling factor α from the slider value
2. Retrieves the next input block from the frames array
3. Processes the block with the current scaling factor
4. Applies overlap-add with the previous block's tail
5. Clips and writes the output chunk to the audio stream
6. Stores the tail portion for the next iteration
7. Cycles through frames in a circular manner for continuous playback using the modulo operator

```

1 prev_tail = np.zeros(BLOCKLEN - Hop)
2 frames = frames_to_process_with_hops(all_samples, BLOCKLEN, OVERLAP_FACTOR)
3 frame_idx = 0
4 while CONTINUE:
5     root.update()
6
7     # get slider value in real time
8     alpha_from_slider = alpha.get()
9
10    # get next input frame
11    input_block = frames[frame_idx]
12
13    # process the scaled
14    output_block = process_block_fft_scaling(input_block, alpha_from_slider)
15
16    # overlap and add

```

```

17     output_block[:len(prev_tail)] = 0.5 * (output_block[:len(prev_tail)] +
18     prev_tail)
19
20     # clipping
21     output_chunk = np.clip(output_block[:Hop], -MAXVALUE, MAXVALUE)
22     output_chunk = np.around(output_chunk).astype(np.int16)
23     stream.write(output_chunk.tobytes())
24
25     prev_tail = output_block[Hop:]
26     # increment to the next frame (circular)
27     frame_idx = (frame_idx + 1) % len(frames)

```

Snippet 4: Main real-time processing loop

- **Parameters:** For the solution, I implemented the following main parameters after a bit of trial and error:

- Block size: 1024 samples
- Overlap factor: 0.5
- Scaling factor range: 0.5 to 2.0

Addendum

Here's the full implementation in code

```
1 import pyaudio
2 import wave
3 import numpy as np
4 import os
5 import math
6 import tkinter as Tk
7
8 # function to make a preprocessed list of frames for overlapped block processing
9 def frames_to_process_with_hops(all_frames, block_size, overlap_factor):
10     frames = []
11     hop_count = math.floor(block_size * (1 - overlap_factor))
12     idx_i = 0
13     idx_l = block_size
14     while idx_l < len(all_frames) - 1:
15         frames.append(all_frames[idx_i:idx_l])
16         idx_i += hop_count
17         idx_l = idx_i + block_size
18     return frames
19
20 # function for processing blocks with scaling via fft and ifft with
21 # interpolation
22 def process_block_fft_scaling(input_block, alpha):
23     X = np.fft.rfft(input_block)
24     Y = np.zeros_like(X)
25     for src_ind in range(X.size):
26         dst_ind = src_ind / alpha
27         if dst_ind < X.size - 1:
28             i0 = int(np.floor(dst_ind))
29             i1 = i0 + 1
30             t = dst_ind - i0
31             Y[src_ind] = (1 - t) * X[i0] + t * X[i1]
32     y = np.fft.irfft(Y)
33     return y
34
35
36 base_dir = os.path.dirname(os.path.abspath(__file__))
37 wavfile = os.path.join(base_dir, 'author.wav')
38 wf = wave.open(wavfile, 'rb')
39 #output_wavfile = 'author_output_blocks_corrected.wav'
40
41 #print('Play the wave file %s.' % wavfile)
42
43 # Open wave file (should be mono channel)
44 #wf = wave.open( wavfile, 'rb' )
45
46 CONTINUE = True # Variable for the looping mechanic
47 CHANNELS      = wf.getnchannels()
48 RATE          = wf.getframerate()
49 WIDTH         = wf.getsampwidth()
50 signal_length = wf.getnframes()
51 BLOCKLEN      = 1024
52 OVERLAP_FACTOR = 0.5
53 MAXVALUE      = 2**15 - 1
```

```

55 print('The file has %d channel(s).' % CHANNELS)
56 print('The frame rate is %d frames/second.' % RATE)
57 print('The file has %d frames.' % signal_length)
58 print('There are %d bytes per sample.' % WIDTH)
59
60 #output_wf = wave.open(output_wavfile, 'w') # wave file
61 #output_wf.setframerate(RATE)
62 #output_wf.setsampwidth(WIDTH)
63 #output_wf.setnchannels(CHANNELS)
64
65 Hop = int(BLOCKLEN * (1 - OVERLAP_FACTOR))
66 binary_data = wf.readframes(signal_length)
67 all_samples = np.frombuffer(binary_data, dtype=np.int16)
68
69 root = Tk.Tk()
70 root.title('Real-time Frequency Scaling')
71
72 # Scaling factor init
73 alpha = Tk.DoubleVar()
74 alpha.set(1.0)
75 # print(alpha.get())
76
77 # Slider config here
78 alpha_slider = Tk.Scale(root, label='Scaling Factor (From 0.5 to 1)', variable=alpha, from_=0.5, to=2.0, resolution=0.01, orient=Tk.HORIZONTAL, length=300)
79 alpha_slider.pack(side=Tk.TOP)
80
81 # Quit button config here
82 def handle_close_quit():
83     global CONTINUE
84     CONTINUE = False
85 B_quit = Tk.Button(root, text='Quit', command=handle_close_quit)
86 B_quit.pack(side=Tk.BOTTOM, fill=Tk.X)
87
88 # Pyaudio config
89 p = pyaudio.PyAudio()
90 # Open audio stream
91 stream = p.open(
92     format      = p.get_format_from_width(WIDTH),
93     channels    = CHANNELS,
94     rate        = RATE,
95     input       = False,
96     output      = True )
97
98
99 # Main loop
100 print('* Start')
101
102 prev_tail = np.zeros(BLOCKLEN - Hop)
103 frames = frames_to_process_with_hops(all_samples, BLOCKLEN, OVERLAP_FACTOR)
104 frame_idx = 0
105 while CONTINUE:
106     root.update()
107
108     # get slider value in real time
109     alpha_from_slider = alpha.get()
110
111     # get next input frame
112     input_block = frames[frame_idx]

```

```

113
114     # process the scaled
115     output_block = process_block_fft_scaling(input_block, alpha_from_slider)
116
117     # overlap and add
118     output_block[:len(prev_tail)] = 0.5 * (output_block[:len(prev_tail)] +
119     prev_tail)
120
121     # clipping
122     output_chunk = np.clip(output_block[:Hop], -MAXVALUE, MAXVALUE)
123     output_chunk = np.around(output_chunk).astype(np.int16)
124     stream.write(output_chunk.tobytes())
125
126     prev_tail = output_block[Hop:]
127     # increment to the next frame (circular)
128     frame_idx = (frame_idx + 1) % len(frames)
129
130
131
132     stream.stop_stream()
133     stream.close()
134     p.terminate()
135     wf.close()

```